

w4112p1

Cody De La Vara, Evan Drewry

March 27, 2013

## **Abstract**

## **Introduction**

## **Main Memory Databases**

## **Column-store Databases**

## **NoSQL and MapReduce**

Developed initially to cope with huge volumes of data that traditional relational database systems have trouble handling efficiently

## **NoSQL**

- NoSQL = "Not only" SQL
- Departure from relational model
- Lightweight and scalable
- Sacrifice consistency for scalability
- Two main types, document stores (Mongo, ...) and key-value stores (BigTable, Cassandra, ...)

## **Document Stores**

### **MongoDB**

### **CouchDB**

## **Key-Value Stores**

### **Cassandra**

## **BigTable**

## **MapReduce**

MapReduce is a simple high-level programming model for processing huge quantities of data in parallel. It is powerful because it provides a layer of abstraction over all the complexities of parallelization—including all partitioning of data among the cluster, execution scheduling, handling of disk and machine failures, and communication between machines—while still providing a simple and flexible programming model.[1]

The

MapReduce is also the name of Google's widely mimicked implementation, however the most popular implementation is Apache's open source Hadoop.

## **Hive**

## **Compared to traditional relational databases**

## References

- [1] Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.